

CFL Alias Analysis

Jia Chen

August 21, 2016

1 Introduction

Alias analysis is one of the most important enabling techniques in an optimizing compiler: without knowing that two pointers cannot alias, performance-critical compiler optimizations like loop invariant code motion, load/store code motion, and dead code elimination could not be applied.

Although LLVM uses an IR of partial-SSA form to trivialize the alias analysis of non-address-taken variables, it still relies on dedicated alias analysis passes to obtain aliasing information for non-promotable memory locations [1]. Existing alias analyses in LLVM codebase (as of May 2016) suffers from various issues:

- Some analyses, such as `basicaa` and `scev-aa`, are intraprocedural and thus cannot look beyond function boundaries.
- Some analyses are very limited in its scope. For example, `globals-aa` only handles non-address take global variables and simple function invocations, while `scev-aa` only handles aliases in loops.
- Some analyses, such as `tbaa`, rely on the existence of IR metadata to detect aliases. If the frontend is not able to emit high-quality metadata, those analyses will not be able to produce many useful results.

To overcome some of those limitations, a new pass called `cfl-aa` was added to LLVM trunk in 2015. The original design goal is to have a complement pass alongside `basicaa` that works on all kinds of pointers, does not depend on metadata, and is fast, more precise, and interprocedural. Before May 2016, some of those goals were achieved, while some useful features were still missing. Also, due to precision and stability issues, `cfl-aa` was never enabled by default in LLVM and clang's production release.

In this year's Google Summer of Code, I worked on a project that aims to bring `cfl-aa` to a usable state. Here is a list of tasks I managed to get done in the past three months:

- Patched various miscompilation bugs in the original `cfl-aa`.
- Improved analysis precision when handling pointer arithmetics and integer casts.
- Added a summary-based interprocedure support to `cfl-aa`.
- Added an experimental inclusion-based `cfl-aa` pass as a more expensive but more precise alternative to the current unification-based `cfl-aa`.
- Evaluated how well `cfl-aa` performs on real-world benchmarks.
- Reported 4 different bugs found on other LLVM components.

In this article, I will try to explain in more detail what I have done in the past three months, and where to go next. It also demonstrates the current design and implementation, and serves as a documentation for those who are interested in working on `cfl-aa` in the future.

The rest of this article is organized as follows: background knowledge required to understand how `cfl-aa` works is provided in Section 2. Section 3 focuses on how the ideas described in Section 2 are implemented. Evaluation of the current implementation can be found in Section 4. Section 5 mainly discuss the limitations of the current system and provide suggestions on how to fix them. Section 6 concludes the article.

2 Background

2.1 Program Representation

We start the discussion by formulating the problem of alias analysis. The input to the analysis is a program represented in LLVM IR form. For the purpose of alias analysis, only values of pointer types are of interest to us¹. In addition, existing alias

¹For simplicity, pointer-to-integer and integer-to-pointer casts are ignored for the rest of the section. We will get back to this issue in Section 3.4.

analyses in LLVM (including cfl-aa) are flow-insensitive, which means that from the analyses' perspective, control flows are simply ignored, and instructions in a function may be executed in any order any number of times.

To simplify the discussion, we assume that any input LLVM IR given to the alias analysis gets canonicalized into the following form:

$$\begin{aligned}
Module &::= Function^* \\
Function &::= Instruction^* \\
Instruction &::= Pointer = \mathbf{alloc} \\
&| Pointer = Pointer' \\
&| Pointer = \mathbf{load} \, Pointer' \\
&| \mathbf{store} \, Pointer, Pointer' \\
&| Pointer = \mathbf{call} \, Function \, (Variable^*) \\
&| \mathbf{return} \, Pointer
\end{aligned}$$

An instruction can be a memory allocation (including stack allocation as well as heap allocation), a pointer assignment, a pointer load/store, or a call/return. Any LLVM instruction can be either put into one of the six forms, or safely ignored.

2.2 Alias Analysis and CFL Reachability

The goal of the alias analysis is to compute alias relations between two *Pointers*. In this article we primarily focus on may-analyses, meaning that given two pointers p and q , an alias analysis only need to produce one of two possible responses: either p and q may alias each other, or they must not alias each other.

A number of approaches have been proposed to solve may-alias analysis problem in the past two decades. The approach adopted by cfl-aa formulates the problem as a context-free language (CFL) reachability problem. Given a canonicalized input program, we first translate it into a graph datastructure called Program Expression Graph [5] (PEG) in the following way:

- Each *Pointer* becomes a node in the graph.
- For each assignment instruction $x = y$, add an edge from node y to node x with edge label **A** to the graph (**A** stands for “assignment”).
- For each load instruction $x = \mathbf{load} \, y$, add an edge from node y to node x with edge label **D** to the graph (**D** stands for “dereference”).
- For each store instruction $\mathbf{store} \, x, y$, add an edge from node y to node x with edge label **D** to the graph.
- How call and return instructions are handled depends on the context sensitivity of the analysis. A simple context-insensitive scheme is to add an **A**-edge from each actual parameter to each formal parameter, and add an **A**-edge from each return value to the left-hand side of each callsite. We will discuss a more precise scheme later in Section 3.5.
- Finally, for each **A**-edge from node x to node y , add another edge from y to x with label $\overline{\mathbf{A}}$, and for each **D**-edge from x to y add another edge from y to x with label $\overline{\mathbf{D}}$.

After the PEG gets constructed, the problem of finding wheter two pointers x and y are aliases then becomes finding whether there is a path in the PEG from the node that corresponds to y to the node that corresponds to x , under the condition that the string formed by concatenating all edge labels on the path can be recognized by the context-free grammar shown in figure 1 [9].

$$\begin{aligned}
\mathbf{V} &::= \overline{\mathbf{F}} \, \mathbf{M?} \, \mathbf{F} \\
\mathbf{M} &::= \overline{\mathbf{D}} \, \mathbf{V} \, \mathbf{D} \\
\mathbf{F} &::= (\mathbf{A} \, \mathbf{M?})^* \\
\overline{\mathbf{F}} &::= (\mathbf{M?} \, \overline{\mathbf{A}})^*
\end{aligned}$$

Figure 1: Context-free grammar for alias analysis

3 Implementation

3.1 Overview

There are two cfl-aa passes that perform CFL Reachability based alias analysis: one is called cfl-steens-aa, and the other one is cfl-anders-aa. They are mostly identical but differs in how the CFL Reachability gets computed. The primary reason for the separation between cfl-steens-aa and cfl-anders-aa is that they each have their own tradeoff between performance and precision. Details about their differences can be found in Section 3.2 and Section 3.3.

Both cfl-aa passes analyze every function separately. Within the function they run in an eager manner: complete aliasing information for the entire function gets computed upfront, to ensure that subsequent alias queries can be answered cheaply. Across function boundaries the analyses adopt a demand-driven strategy. Only when a client makes alias queries on a function do we compute the aliasing information for that function.

The canonicalization procedure described in Section 2.2 is implemented in a class called *CFLGraphBuilder*. Source code for the class is located at *lib/Analysis/CFLGraph.h*. Both cfl-steens-aa and cfl-anders-aa rely on it to create the aforementioned PEG datastructure (which is called *CFLGraph* in our implementation).

3.2 Steensgaard-Style CFL Reachability

After *CFLGraph* is constructed, the next task is to compute pairwise CFL Reachability relations for all pointers in the given function. This is where cfl-steens-aa and cfl-anders-aa start to look different from one another. In this section we describe the algorithm used by cfl-steens-aa first², and the one used by cfl-anders-aa is the focus of the next section.

It is known in the literature that the problem of computing general CFL Reachability can be reduced to the problem of computing a dynamic transitive closure, which is not known to have subcubic solutions [5]. Even if we only consider the CFL Reachability problem with a fixed grammar shown in figure 1, I am still not aware of any algorithm that can do better than $O(n^2)$, which, in some cases, is too expensive to be added LLVM’s pass pipeline.

The cfl-steens-aa pass utilizes on a simple observation to drastically improve the worst-case complexity bound: in figure 1, if we ignore nonterminal **A** and nonterminal $\bar{\mathbf{A}}$, then the remaining grammar essentially describes a *Dyck* language (i.e. a language that only recognized balanced parentheses). If the underlying context-free language is Dyck, CFL Reachability can be computed in linear (if the corresponding graph is a tree) or log-linear (if the corresponding graph is more general than a tree) time, with the right choice of data structure [8].

The way cfl-steens-aa ignores **A** and $\bar{\mathbf{A}}$ is to perform node merging on *CFLGraph*: if two nodes on the graph are connected by an **A**-edge or an $\bar{\mathbf{A}}$ -edge, the two will be collapsed into a single node. We repeat this process until all **A** and $\bar{\mathbf{A}}$ edges are removed from the graph. Then we apply the $O(n)$ algorithm proposed in [8] to compute the Dyck CFL Reachability of the simplified graph. The algorithm will eventually group pointers into equivalent sets, and we can tell if two pointers may-alias each other by checking whether or not they are in the same equivalent set.

The node collapse trick mentioned previously essentially approximate nonterminal **V** and nonterminal **M** in the original grammar by their transitive closures. It overapproximates the CFL Reachability of the original grammar, meaning that the analysis could answer *MayAlias* even if in the original grammar the answer should have been *NoAlias*. In terms of analysis precision, it has been shown in [9] that treating **V** and **M** as transitive relations lead to an analysis that is equivalent to Steensgaard’s unification-based algorithm [7].

3.3 Andersen-Style CFL Reachability

Although cfl-steens-aa successfully brings the cost of CFL Reachability computation down, it achieves the goal by sacrificing analysis precision. Consider the following code snippet containing three pointers *x*, *y* and *z*:

```
x = alloc;  
y = alloc;  
z = x;  
z = y;
```

When running cfl-steens-aa on it, all three pointers gets collapsed into the same node, which means that cfl-steens-aa cannot distinguish them from one another: the three pointers are treated as pairwise aliases by the analysis. However, it is obvious from the code that the two pointers *x* and *y* should not alias each other, since they each point to a different memory allocation.

Here is another example:

```
x = alloc;  
y = alloc;
```

²Note that I take no credit from what gets demonstrated in this section. The interprocedural version of cfl-steens-aa in LLVM was 90% finished before I started the project. Nevertheless, other components described in this section, such as the cfl-anders-aa pass and the interprocedural extension, are mostly my own work.

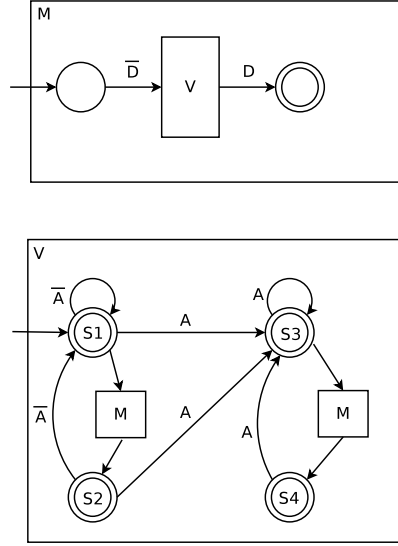


Figure 2: Hierarchical State Machine used by Zheng et al [9].

```
*a = x ;
*b = y ;
*c = x ;
*c = y ;
```

If we examine the snippet manually, we can see that pointer a and b do not alias each other, but both of them alias c . But if we feed this program to `cfl-steens-aa`, it will try to unify both (a, b) and (b, c) , therefore creating spurious alias pair (a, c) .

`Cfl-steens-aa` tries to gain analysis speed at the cost of analysis precision. In contrast, its cousin `cfl-anders-aa` picks another tradeoff, where precision is valued over performance. In `cfl-anders-aa`, the original grammar shown in figure 1 is used. No overapproximation is performed. The justification for building `cfl-anders-aa` is that, although the theoretical bound to compute CFL Reachability is cubic, in practice the worst-case bound is rarely hit [6]. Also, there has been some evidences that suggest with the right set of optimization implemented, analyses that are equivalent to `cfl-anders-aa` can scale well [4][3]. One additional consideration is that it is often easier for non-unification-based algorithms to capitalize on field-sensitivity. And finally, even if `cfl-anders-aa` never gets practical in terms of performance, it is still worth implementing just to see how much more precision we can get compared to `cfl-steens-aa`. If anyone needs the extra precision, they can always turn it on.

The algorithm used for `cfl-anders-aa` comes from existing work [9]. The basic idea of the algorithm is to perform a traditional transitive closure computation to propagate the reachability information, but stop propagating if it is found that the edge labels do not form a string that can be recognized by the CFL grammar. To facilitate CFL grammar recognition, the grammar is encoded using a hierarchical, recursive state machine, and we can tell whether we should propagate CFL reachability along a given edge by checking if the label on the edge can direct us from an accepted state in the state machine to another accepted state. Pseudocode of the algorithm can be found in Figure 4 of the paper.

Despite the similarity, in `cfl-anders-aa` we do deviate from the algorithm described in the paper [9] in the following two ways. First, our algorithm eagerly computes all alias pairs after the `CFLGraph` is built, while in the paper the authors did the computation in a demand-driven fashion. This is achieved by initializing the worklist with all neighbors in `CFLGraph`, instead of only putting the pair being queried into the worklist. We did not implement the demand-driven algorithm due to the additional coding complexity and higher memory profile. Second, in the paper the authors use a state machine that does not distinguish value reads from value writes. We try to differentiate the two, as the information of whether a given pointer is read from or written to turned out to be crucial for building alias summaries as well as modref summaries. Our implementation ends up using a modified version of the paper’s state machine (shown in figure 3), where some states in the original machine (shown in figure 2) were duplicated just to ensure value reads and value writes were not conflated.

It has been shown that the algorithm we use here is equivalent to Andersen’s inclusion-based algorithm [2] in terms of analysis precision.

3.4 Node Attributes

Certain aspects in LLVM IR cannot be properly encoded into `CFLGraph`:

- LLVM allows casts between pointers and integers, but `CFLGraph` only contains pointers as its nodes.
- LLVM module may contain global variables, which do not belong to any function and thus not modeled by `CFLGraph`.

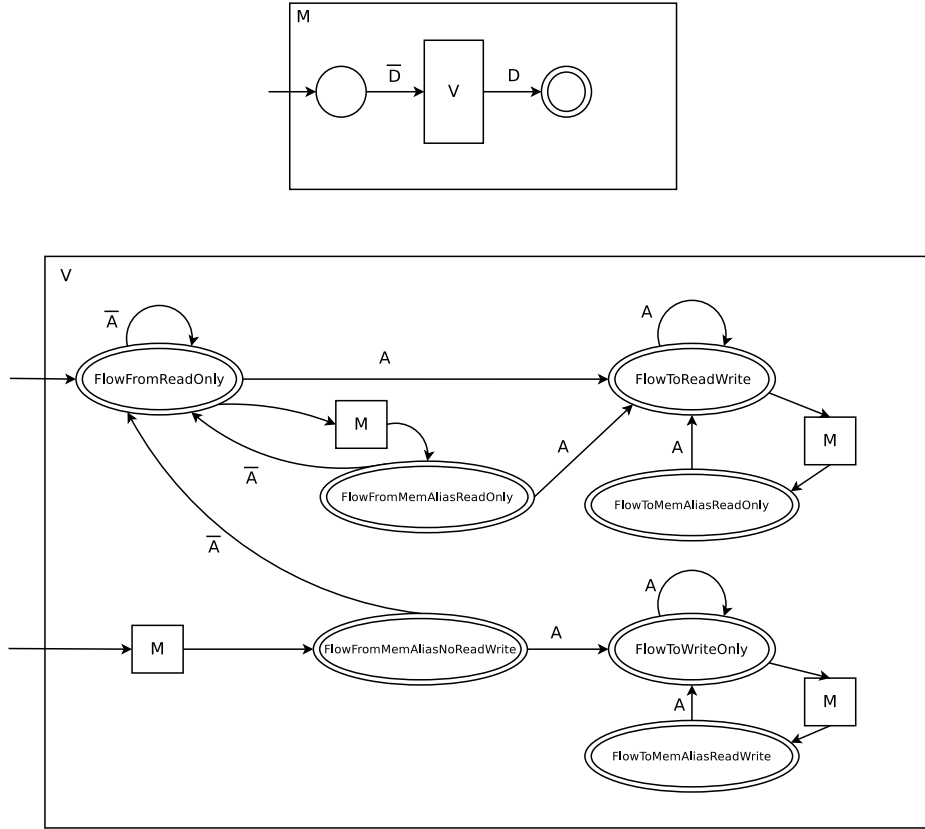


Figure 3: Hierarchical State Machine used in cfl-anders-aa.

- LLVM module may contain declared but not defined functions. CFLGraph has no way of knowing what their side-effects they have and therefore must assume that they can write anything to any memory location they get access to.

To cope with the aforementioned language features, each node in CFLGraph is associated with zero or more attributes. Attributes serve as “escape hatches” to our CFL Reachability algorithm: when a client queries the alias relation between pointer a and b , if either of them gets tagged with one or more attributes, the alias analysis will skip looking up their CFL reachability lookup and yield a response directly according to figure 3.4. Attributes also has the property that they get propagated “downwards”: if a gets tagged with certain attributes, then everything that is transitively pointed to by a also gets tagged with the same set of attributes.

There are five different attributes: *AttrEscaped*, *AttrCaller*, *AttrGlobal*, *AttrArgument*, and *AttrUnknown*. Here is a brief explanation of their meaning:

- Pointers that are allocated locally but somehow assigned to another value not modeled by the analysis get tagged with the *AttrEscaped* attributes. Pointers that are casted into integers, or pointers that are used as actual parameters of some opaque function call, all fall into this category.
- Formal parameters of the current function get tagged with *AttrArgument*. Values pointed to by those parameters get tagged with *AttrCaller*.
- Global values themselves get tagged with *AttrGlobal*.
- Any values that come from an unresolvable source (e.g. fabricated from integers, or loaded from opaque memory) are all tagged with *AttrUnknown*.

3.5 Interprocedural Analysis

The algorithm described in Section 3.2 and Section 3.3 both work intraprocedurally. This section demonstrates how to extend those algorithms to work in an interprocedural setting.

In general, there are two distinct approach to perform interprocedural analysis: the top-down approach and the bottom-up approach. A top-down approach analyzes a function *before* it finishes analyzing all of its callee, while bottom-up approach (a.k.a. summary-based approach) analyzes a function *only after* it finishes analyzing all of its callees. Top-down approaches are generally more precise than bottom-up approaches, since when analyzing the callees the analysis can put them in a known

	No Attributes	AttrEscaped	AttrGlobal	AttrArgument	AttrCaller	AttrUnknown
No Attributes	CFL lookup	CFL lookup	NoAlias	NoAlias	NoAlias	NoAlias
AttrEscaped	CFL lookup	CFL lookup	NoAlias	NoAlias	MayAlias	MayAlias
AttrGlobal	NoAlias	NoAlias	MayAlias	MayAlias	MayAlias	MayAlias
AttrArgument	NoAlias	NoAlias	MayAlias	MayAlias	MayAlias	MayAlias
AttrCaller	NoAlias	MayAlias	MayAlias	MayAlias	MayAlias	MayAlias
AttrUnknown	NoAlias	MayAlias	MayAlias	MayAlias	MayAlias	MayAlias

Figure 4: Rules for attribute handling. Rows represent the attributes of one pointer, and columns represent the attributes of another pointer. Each cell represents what response the analysis should return according to the attributes of the two given pointers. If there are more than one attributes for each pointer, combinations that lead to *MayAlias* are prioritized over those that lead to *NoAlias*.

context. In contrast, in a bottom-up approach functions must be examined independently and out-of-context, where we have zero information on the state of the memory, the arguments, the global values, or anything else external to the function. To carry out the analysis conservative assumptions have to be made about those external states. In exchange for the potential loss of precision, we get better scalability: each function only needs to be analyzed once, and the analysis result we obtain for each function is highly reusable and can be applied to anywhere the function is called.

Both cfl-aa passes use a bottom-up approach to look beyond function boundaries. Each function gets analyzed independently and intraprocedurally at first, and we store the externally visible effects of the function into a *function summary*. The format of those summaries is listed in figure 5. Subsequently, we *instantiate* each summary at each callsite of the corresponding function, replacing **ReturnValue** with the right-hand side of the callsite, replacing **Parameter**(*i*) with the *i*-th actual parameter of the callsite, and tagging the return value and the parameters with the right set of attributes. Finally, we incorporate the instantiated summary into the caller’s CFLGraph, and now the analysis can proceed as usual. In other words, the interprocedural part of the analysis is handled entirely at the CFLGraph construction phase. The CFL Reachability computation is completely oblivious of the function summaries, and therefore we do not need to make any change to it for the interprocedural analysis to work.

$$\begin{aligned}
\text{FunctionSummary} &::= \text{ExternallyVisibleEffect}^* \\
\text{ExternallyVisibleEffect} &::= \text{InterfaceValue} = \text{InterfaceValue}' \\
&| \text{InterfaceValue} \leftarrow \text{Attribute} \\
\text{InterfaceValue} &::= \mathbf{ReturnValue} \\
&| \mathbf{Parameter}(i) \\
&| * \text{InterfaceValue}' \\
\text{Attribute} &::= \mathbf{AttrEscaped} \\
&| \mathbf{AttrGlobal} \\
&| \mathbf{AttrUnknown}
\end{aligned}$$

Figure 5: Syntax of function summary

4 Evaluation

In this section, we evaluate both cfl-steens-aa and cfl-anders-aa using all C/C++ programs in SPEC2006 benchmark. Source codes of the test programs were first compiled into LLVM bitcode file separately with optimization level 0, and then linked together into a single bitcode file before any analysis gets performed. The experiments were performed on a Ubuntu 14.04 environment, with Core i5 4330M as the CPU and 8GB of RAM. Native gcc toolchain version on the system was 5.3.0. The llvm and clang trunk revision used for the experiments is r278855.

4.1 Alias Responses

In the first set of experiments, we use the **-aa-eval** pass to query the pairwise alias relations in the input bitcode, and count the percentage of *MayAlias* responses we get from each alias analysis under test. More precise analyses tend to generate less *MayAlias* responses, so lower percentage is better. We test five different alias analysis settings:




Figure 6: Percentage of may-alias responses (running on unoptimized bitcodes)

- Using basicaa alone
- Using cfl-steens-aa alone (referred to as “steens” for short)
- Using cfl-anders-aa alone (referred to as “anders” for short)
- Using basicaa combined with cfl-steens-aa (referred to as “basicsteens” for short)
- Using basicaa combined with cfl-anders-aa (referred to as “basicanders” for short)

The result can be found in figure 6³. For most of the benchmarks, we can see that running cfl-steens-aa or cfl-anders-aa alone yields worse result than running basicaa alone. One explanation is that basicaa has been taught to handle far more cases than those cfl-aas do, including queries that contain global values and `getelementptrs`. As a result, both cfl-aas should not be used as a replacement of basicaa.

Putting cfl-aas behind basicaa, as figure 6 suggests, behaves consistently better than running basicaa alone. In some of the benchmarks (e.g. `hmmmer` and `omnetpp`), the precision gain is rather significant. Benchmarks on which the precision gains are small (e.g. `bzip2` and `sjeng`) tend to use structs and arrays extensively, and currently neither version of cfl-aa can handle them in a precise manner.

We also performed the same set of experiments after running `-mem2reg` on the input program, as `-mem2reg` is able to remove a large amount of unnecessary stack allocation and can reveal the program’s memory-related behaviors. Results are shown in figure 7. Under this scenerio, the precision gap between basicaa and both versions of cfl-aas becomes even bigger, but we do see that there are still times when running basicaa together with cfl-aas could give basicaa a large precision boost.

Figure 8 and figure 9 present the running time for the two sets of experiments (lower is better). Data were collected using a release build of `llvm` and `clang`. From the plot we can see that if we take basicaa as a baseline, standalone cfl-steens-aa almost always runs significantly faster, and standalone cfl-anders-aa usually runs noticeably faster. It shows that basicaa is indeed a highly sophisticated analysis that can do much more than both cfl-aas could handle⁴. When we put cfl-steens-aa behind basicaa, normally we do not see much performance degradation except for a few benchmarks. Running cfl-anders-aa with basicaa lead to a big slowdown, and if we also consider the precision numbers from figure 6 and figure 7, this combination is simply not worth not worth to use.

³Cfl-anders-aa did not terminate for `perlbench` and `gcc` in 30 minutes. As a result corresponding data were not shown in the graph.

⁴However, one important thing to point out here is that basicaa is a demand-driven analysis, and the performance numbers here were obtained by exhaustive alias pair queries. In practice, since clients rarely behave this way, we would expect the cost of basicaa to go down while the cost of both cfl-aas remains.




Figure 7: Percentage of may-alias responses (running after **-mem2reg**)




Figure 8: Analysis time (running on unoptimized bitcodes)




Figure 9: Analysis time (running after **-mem2reg**)

4.2 Impact to the clients

In this section we evaluate cfl-aa by feeding their results to two optimization passes: loop invariant code motion and global value numbering. Compared with the **-aa-eval** pass, licm and gvn are more realistic clients and therefore can give us a better idea of cfl-aa’s utility in practice. All statistics here were collected by building LLVM with assertions, adding **-stats** to the command line arguments, and then parsing the output.

The first client we investigated is the licm pass. Measurement collected for this client is the number of instructions hoisted out of the loop. Results are shown in table 1. The second client we picked is the gvn pass⁵. Here we collected the number of instructions removed by gvn, plus the number of instructions PRE’d by gvn, as both numbers contributes to the quality of the optimization. The gvn results are shown in table 2

We observe two trends in the dataset: first, cfl-steens-aa and cfl-anders-aa, when used along with basicaa, consistently help to improve the effectiveness of the clients. Second, when comparing the column of *basicsteens* and *basicanders*, we do not see improvement at similar order of magnitude.

5 Limitation

A properly functional alias analysis pass, just like many other analysis passes, requires a large amount of tuning to work properly. Cfl-steens-aa and cfl-anders-aa are no exception to this rule. In this section, we try to summarize the shortcomings of the existing solution, and suggest ways to patch or alleviate those problems in the future.

Cfl-anders-aa is slow As shown in the evaluation section, cfl-anders-aa takes a long time to run, and in some cases it even fails to terminate within a reasonable amount of time. This is due to the lack of bug hunting and performance tuning: compared with cfl-steens-aa, which has been curated by various developers for almost a year, cfl-anders-aa was a one-man work created weeks ago. As an inclusion-based context-sensitive analysis, cfl-anders-aa inherently has higher computation cost, and therefore it heavily relies on various form of optimizations [4][3], as well as performing demand-driven style calculations [9], to boost its scalability. Currently none of those optimizations get implemented, as at this early stage we choosed to prioritize correctness over performance. Hopefully as cfl-anders-aa matures, those optimizations can be added on top of the existing prototype.

⁵Technically speaking, gvn does not use `AliasAnalysis` directly. But it heavily relies on `MemoryDependenceAnalysis`, which in turn depends on `AliasAnalysis`.

Table 1: Numbers of instructions hoisted in licm (running on unoptimized bitcodes)

	basic	steens	anders	basicsteens	basicanders
400.perlbench	3791	915	-	3969	-
401.bzip2	1096	2546	2673	2901	3010
403.gcc	19145	3739	-	19788	-
429.mcf	169	97	150	217	217
433.milc	4811	2743	2830	4833	4833
444.namd	8511	3185	4390	8713	9287
445.gobmk	9049	4557	4597	10620	10634
447.dealII	48472	21577	21624	48484	48484
450.soplex	7618	5166	5472	7642	7644
453.povray	11822	6289	6508	12061	12153
456.hmmer	5911	5158	5783	9400	9923
458.sjeng	969	637	637	1109	1123
462.libquantum	1204	528	528	1228	1228
464.h264ref	17589	19089	20978	28267	29563
470.lbm	148	432	467	520	520
471.omnetpp	1947	975	984	1954	1954
473.astar	926	613	661	930	978
482.sphinx3	3758	1637	1819	4156	4191
483.xalancbmk	14631	8832	9071	14652	14658

Table 2: Number of instructions removed or PRE'd in gvn (running on unoptimized bitcodes)

	basic	steens	anders	basicsteens	basicanders
400.perlbench	100135	79323	-	100861	-
401.bzip2	9465	7931	9062	9864	9869
403.gcc	300212	249682	-	301410	-
429.mcf	1431	957	1387	1456	1467
433.milc	13040	11990	12531	13232	13267
444.namd	34181	31494	34510	35026	35030
445.gobmk	75309	71925	72263	76717	76721
447.dealII	165369	161826	164245	165634	165642
450.soplex	26594	25786	27281	27206	27715
453.povray	72402	66833	69825	73430	73487
456.hmmer	30951	32796	34440	35117	35963
458.sjeng	9734	7552	7610	9903	9908
462.libquantum	3299	3123	3123	3284	3284
464.h264ref	64530	55566	59183	68659	69545
470.lbm	2335	2712	2727	2730	2730
471.omnetpp	22856	21409	22166	22859	22860
473.astar	3626	3334	3491	3715	3789
482.sphinx3	15621	14276	15553	16190	16302
483.xalancbmk	193876	184509	191352	194205	194289

Cfl-aa is field-insensitive Currently both versions of cfl-aa treat structs as a giant data blob, and do not try to differentiate between different fields inside the same struct. However, it is usually the case that if a struct contains more than one pointer field, those fields tend to point to different locations. Collapsing structs causes a great loss in terms of precision, and it may be one of the reasons why standalone cfl-aas get vastly outperformed by basicaa in terms of precision in the evaluation section.

There has been some effort spent on adding field sensitivity to cfl-anders-aa: CFLGraph has been made field-offset aware, and the CFL Reachability propagation logic already gets the ability to perform offset arithmetics. However, to support field sensitivity at the interprocedural level, we need to adjust the format of alias summaries and incorporate field offset as part of *ExternallyVisibleEffect*. This is the only issue left to be addressed before we can put cfl-anders-aa into a full field-sensitive mode.

On the other hand, no work was done to add field sensitivity support to cfl-steens-aa. The decision to prioritize field-sensitive cfl-anders-aa over field-sensitive cfl-steens-aa was made under the original hypothesis that field-sensitivity may be less helpful to cfl-steens-aa due to its unification-based nature. Given the evaluation numbers, it might be worth reconsidering the validity of the hypothesis and add field-sensitive cfl-steens-aa into our short-term goal.

Interprocedural analysis is imprecise for recursive calls In Section 3.5, we mentioned that both cfl-aa passes analyze functions in a bottom-up order: callees are examined before the callers. But if we have potential recursive calls in the program (i.e. strongly connected components in the call graph), a caller may be the callee of itself, on which case we need additional rules to define the analysis order.

The current solution to break the recursive cycle is to pick a function on the loop and conservatively treat it as an externally defined function (i.e. may write anything to any memory location it has access to). Which function gets picked depends on which function gets analyzed first, which in turn depends on what function the client queries the alias analysis first.

Our current approach is sound, but may be imprecise since there is no guarantee that precision-critical functions will not be treated conservatively. A more precise approach is to start the analysis with the assumption that all functions have no side effect, then as we discover more side effects by analyzing the body of each function, we propagate the effects back to its callers. If the propagation adds even more side effects to the caller, propagate them to the callers of the caller, and so on. We repeat the entire process until a fixpoint is reached. This fixpoint strategy may be more costly than our current approach, since we no longer have the guarantee that each function gets analyzed only once: when there are cycles on the call graph, functions on the call graph may be examined more than once due to side effect propagation. Whether the precision gain is worth the performance cost is another question to explore.

Cfl-aa lacks support for modref queries The `AliasAnalysis` interface in LLVM contains a wide range of APIs. Both cfl-anders-aa and cfl-steens-aa only support the `alias()` queries. However, `alias()` is not the only way clients interact with `AliasAnalysis`. Another popular class of `AliasAnalysis` APIs are those `getModRefInfo()` queries. The `getModRefInfo()` methods return information about whether a given callsite (or a given callee) may read from or write into certain memory locations. The information is very important to other optimization passes that touch call or invoke instructions, because without `getModRefInfo()` those passes would, say, refuse to move or delete redundant function invocations, even if the alias analysis is 100% confident that the invocation is free of side effects.

There used to be naive modref supports in both CFL-based analyses: the `getModRefInfo()` methods will just look into function summaries of the alias analysis, and derive modref information accordingly. The support was removed later, as we found that using alias analysis function summaries to derive modref information was unsound: the summaries for alias analysis contain values of pointer-type only, but a sound modref analysis must also consider reads/writes of non-pointer values. If a function accesses a non-pointer value, the fact will not be reflected in the alias summary, but it should be caught by the modref analysis.

A modref analysis may utilize alias summaries to figure out what memory locations are accessed, but conceptually it is a separate analysis that has to consider a wider range of values and therefore requires more work than our naive implementations.

Collaboration with basicaa LLVM's basicaa pass is surprisingly efficient and precise, as shown in the evaluation section. It is inherently demand-driven, handles a wider selection of values, and even has the ability to decompose GEPs — all of these nice features strongly suggest that basicaa is not something that can be easily replaced. The word “basic” in basicaa does not mean “naive” or “elementary”. Instead, it means “essential” and “indispensable”.

If basicaa is irreplaceable, then when it claims that two pointers are not aliases, it might not be a good idea for other CFL-based passes to spend more time just to rediscover the same fact. A more detailed analysis of basicaa's strength and weaknesses may be necessary, to provide better insights into what aspects cfl-aas should be enhanced so that the enhancement does not reimplement what basicaa can already do.

Cfl-aa does not re-execute when function body gets changed Technically speaking, this is less a limitation of cfl-aa itself, and more of a limitation of the LLVM alias analysis framework. However, the problem does have a large impact to the utility of cfl-aas, so I think it is worth to be pointed out here.

We turn on `cfl-anders-aa` or `cfl-steens-aa` by specifying `-use-cfl-aa` on the command line when optimization level is larger than 0, what the pass manager will do is to put `cfl-anders-aa` or `cfl-steens-aa` at the front of the pass pipeline, run it once, and never re-execute it afterwards. What happens then is that the alias summaries generated by `cfl-aas` quickly becomes outdated after several transformation passes take effects. Those transformation passes could potentially add values to or remove values from the function body, and to stay sound `cfl-aas` have to conservatively return *MayAlias* if one or more values in the aliasing queries are not found in its summary. If a large amount of values in the function are not there `cfl-aa` was first executed, then we might not be able to obtain useful responses. As a result, only those transformation passes that are scheduled “early” on the optimization pipeline enjoy the benefit of `cfl-aas`. Passes that are scheduled “late” will behave almost the same, regardless of whether `cfl-aas` are turned on or not.

It would be nice to have the pass manager periodically re-execute `cfl-aas` after certain number of transformation passes, or to allow transformation passes to explicitly invalidate `cfl-aa` summaries such that they could get automatically recomputed by `cfl-aas`.

6 Conclusion

In this article we discussed about the design and implementation of two flow-insensitive, context-sensitive LLVM alias analysis passes. Both passes are based on a CFL Reachability solver with different tradeoff between precision and performance. We evaluated those passes using real-world clients and benchmarks, and proposed various ways to improve upon the current implementations.

The original title of the project is called *Better Alias Analysis By Default*. Based on the evaluation numbers, `basicaa+cfl-steens-aa` might be a promising candidate to set as the default alias analysis for LLVM, if the re-execution issue mentioned in Section 5 can be properly addressed.

Acknowledgement

Many thanks to my mentors George Burgess IV, who helped me get started on the project and reviewed all of my patches, and Hal Finkel, who provided me with access to high-performance platforms as well as valuable information on testing and debugging.

I would also like to show my gratitude to Daniel Berlin for sharing his ideas and knowledge with me through email discussions.

Kudos to all other devs on `phabricator` who have reviewed my codes.

References

- [1] <http://llvm.org/docs/AliasAnalysis.html>.
- [2] Lars Ole Andersen. Program analysis and specialization for the c programming language. Technical report, 1994.
- [3] Ben Hardekopf and Calvin Lin. The ant and the grasshopper: Fast and accurate pointer analysis for millions of lines of code. In *Proceedings of the 28th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI ’07*, pages 290–299, New York, NY, USA, 2007. ACM.
- [4] Ben Hardekopf and Calvin Lin. Exploiting pointer and location equivalence to optimize pointer analysis. In *Proceedings of the 14th International Conference on Static Analysis, SAS’07*, pages 265–280, Berlin, Heidelberg, 2007. Springer-Verlag.
- [5] Thomas Reps. Program analysis via graph reachability. In *Proceedings of the 1997 International Symposium on Logic Programming, ILPS ’97*, pages 5–19, Cambridge, MA, USA, 1997. MIT Press.
- [6] Manu Sridharan and Stephen J. Fink. *The Complexity of Andersen’s Analysis in Practice*, pages 205–221. Springer Berlin Heidelberg, Berlin, Heidelberg, 2009.
- [7] Bjarne Steensgaard. Points-to analysis in almost linear time. In *Proceedings of the 23rd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL ’96*, pages 32–41, New York, NY, USA, 1996. ACM.
- [8] Qirun Zhang, Michael R. Lyu, Hao Yuan, and Zhendong Su. Fast algorithms for dyck-cfl-reachability with applications to alias analysis. In *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI ’13*, pages 435–446, New York, NY, USA, 2013. ACM.
- [9] Xin Zheng and Radu Rugina. Demand-driven alias analysis for c. In *Proceedings of the 35th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL ’08*, pages 197–208, New York, NY, USA, 2008. ACM.